




Facial Emotion Detection

Jeff Stearns

PROBLEM DEFINITION

- Can emotion be accurately detected by Convolutional Neural Networks (CNN)?
- Average natural human ability to detect emotions is around 90%, can we approach this level?



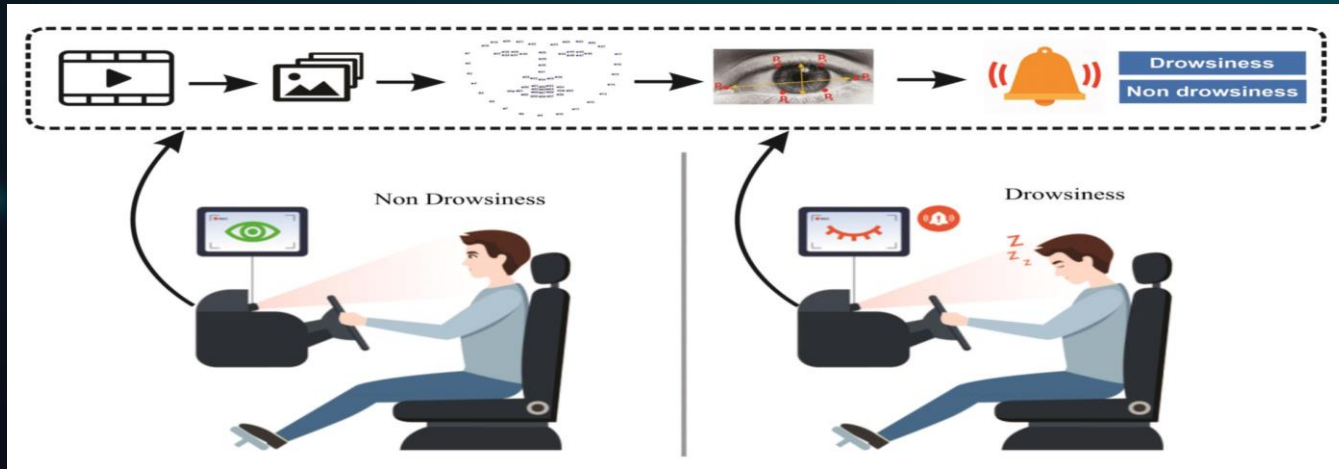
EMOTION DETECTION USE CASES

- Self driving car, detect drowsy or distracted.
- Potentially assist with security: Predict dangerous behaviors by reading emotions.
- Detect emotional response of customers to advertisement.



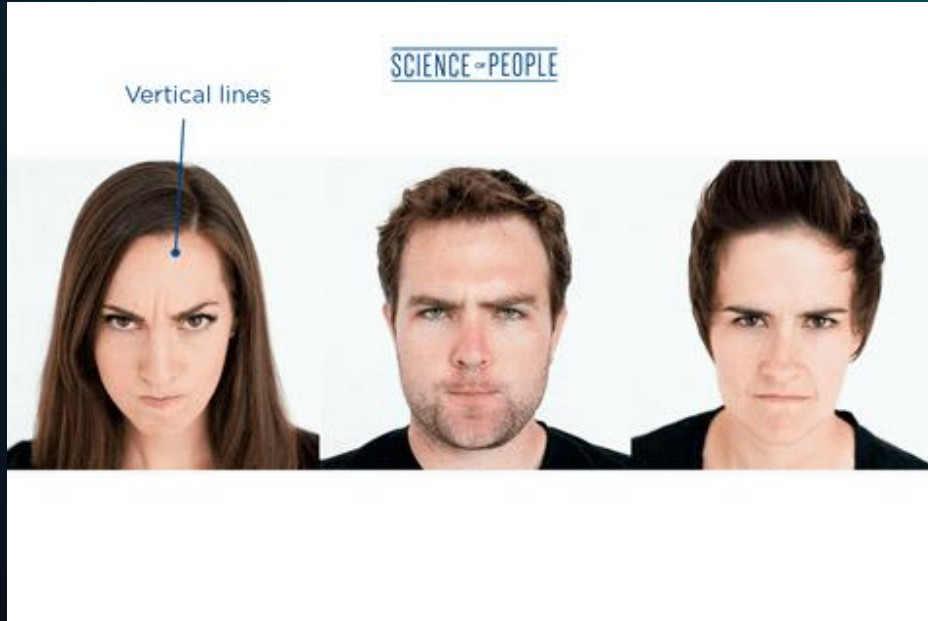
FACIAL EMOTION DETECTION USE CASE 1

Autopilot Automobiles : Driver alertness still required, sensors and software are not perfect.




FACIAL EMOTION DETECTION USE CASE 2

Security alert, detect potentially dangerous individual through emotion recognition.



FACIAL EMOTION DETECTION USE CASE 3

- Detect emotional response of customers to advertisement.



DATA SET ANALYSIS

Facial_emotion_images.zip

- Black and White Images : 48 Pixel Squares
- Example Images

Happy



Neutral



Sad



Surprise



DATA SET ANALYSIS

- 20,214 Grayscale Images, 48 Pixel Square
- Very small TEST set : 60-80 / 10-20 / 10-20 standard ratios

TRAINING	VALIDATION	TEST
Happy 3976 26.3% Sad 3982 26.4% Neutral 3978 26.3% Surprise 3173 21.0%	Happy 3976 36.7% Sad 3982 22.9% Neutral 3978 24.4% Surprise 3173 16.0%	Happy 32 25.0% Sad 32 25.0% Neutral 32 25.0% Surprise 32 25.0%
15109 74.7%	4977 24.6%	128 0.6%

DATA SET ISSUES


No face, watermark, multiple face, cartoon, etc...



train\happy



















train\sad



test\happy

DATA SET IMAGES, NO FACE or TOO DARK

- Filter Tool : Reject Images $\text{Numpy.Average} < 20$ or > 235
- 25 Images : 21 Train, 3 Validation, 1 Test





CONVOLUTIONAL NEURAL NETWORK (CNN) MODELS

MODEL	LAYERS	PARAMETERS	COMMENT
Grayscale 1	5 : 3 Conv2d	605,060	Lowcode Model
Grayscale 2	6 : 3 Conv2d	389,604	BatchNormalization, Dense
VGG16	16 : 13 Conv2d	14,714,688	Transfer Model 1, RGB Images
ResNetV2	164 : 1000 Categories	42,658,176	Transfer Model 2, RGB Images
EfficientNet	B7 : 813	8,769,374	Transfer Model 3, RGB Images
Milestone 1	7 : 4 Conv2d	1,592,324	Kernel Size 3 & 2, Conv2D added
CapStone	8 : 5 Conv2d	2,973,700	Conv2D Layer Added

DATASET LOADERS

TENSORFLOW TOOLS

- ImageDataGenerator : Image Augmentations Applied
 -  • rescale : Normalize pixel values 0-1
 - horizontal_flip : Teach model differing orientations
 - rotation_range : Handle random rotation of images
- flow_from_directory method
 -  • Read data, labels from directory structure
 - Easier to load and process data in batches



CNN MODEL 1

Low Code Notebook Model


- We want our Base Neural Network architecture to have 3 convolutional blocks.
- Each block contains one Conv2D layer, maxpooling layer and Dropout layers.
- Add first Conv2D layer with **64 filters** and a **kernel size of 2**, 'same' padding **input shape = (48, 48, 1) if you're using 'grayscale' colormode**. Use **'relu' activation**.
- Add a second Conv2D layer with **32 filters** and a **kernel size of 2**, 'same' padding and **'relu' activation**.
- Add a third Conv2D layer with **32 filters** and a **kernel size of 2**, 'same' padding and **'relu' activation**.
- After adding your convolutional blocks, add your Flatten layer.
- Add your first Dense layer with **512 neurons**. Use **'relu' activation function**.
- Add a Dropout layer with dropout ratio of 0.4.
- Add your final Dense Layer with 4 neurons and **'softmax' activation function**
- Print your model summary

CNN MODEL 2

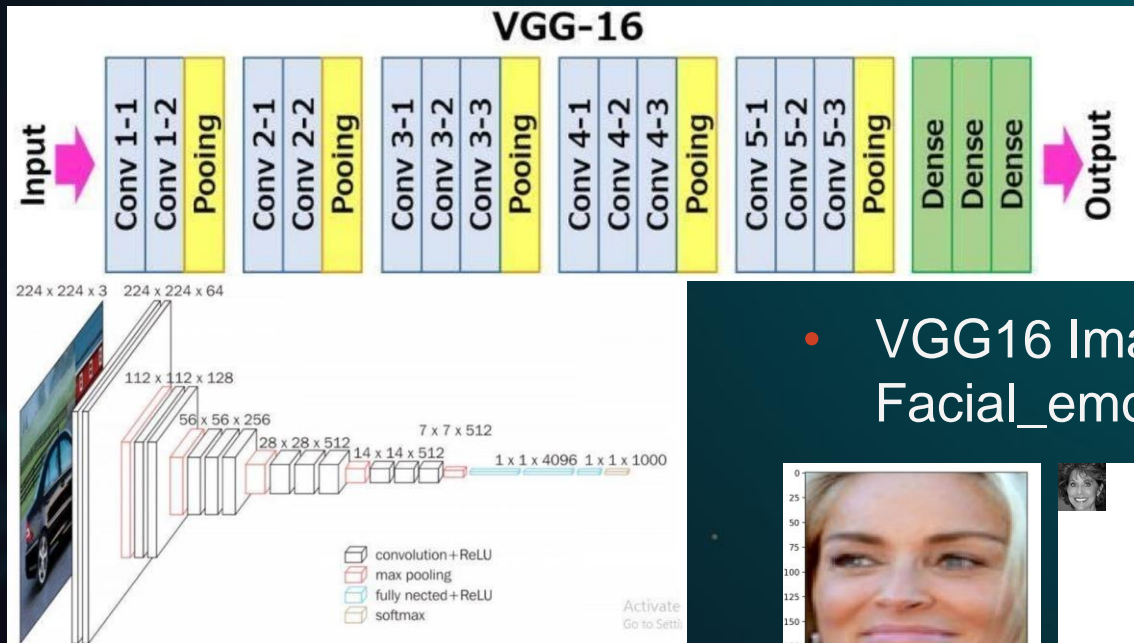
Enhanced Low Code Notebook Model

- Conv2D Layer Added
- Dense Layer Added
- More Neurons per Conv2D Layer
- BatchNormalization Added


CNN MODELS 1 & 2 GRAPH



VGG-16 ARCHITECTURE



- VGG16 Image vs Facial_emotion_images





- 65 to 1 Image Data

ResNetV2 ARCHITECTURE

- 42,658,176 Parameters

Compressed View



- Poor performance on Dataset : ~25%

EFFICIENTNET ARCHITECTURE

- 8,769,374 Parameters




- Poor performance on Dataset : ~25%

MILESTONE 1 MODEL

Enhanced Version 3 CNN Model


- Conv2D Layer Added
- More Neurons per Conv2D Layer
- Data Loader Augmentation Hyperparameter adjustments
- Validation accuracy dips to in Epoch 8?



CAPSTONE MODEL


Final CapStone Model Version

- Conv2D Layer Added
- More Neurons per Conv2D Layer
- More parameters : 2,973,700

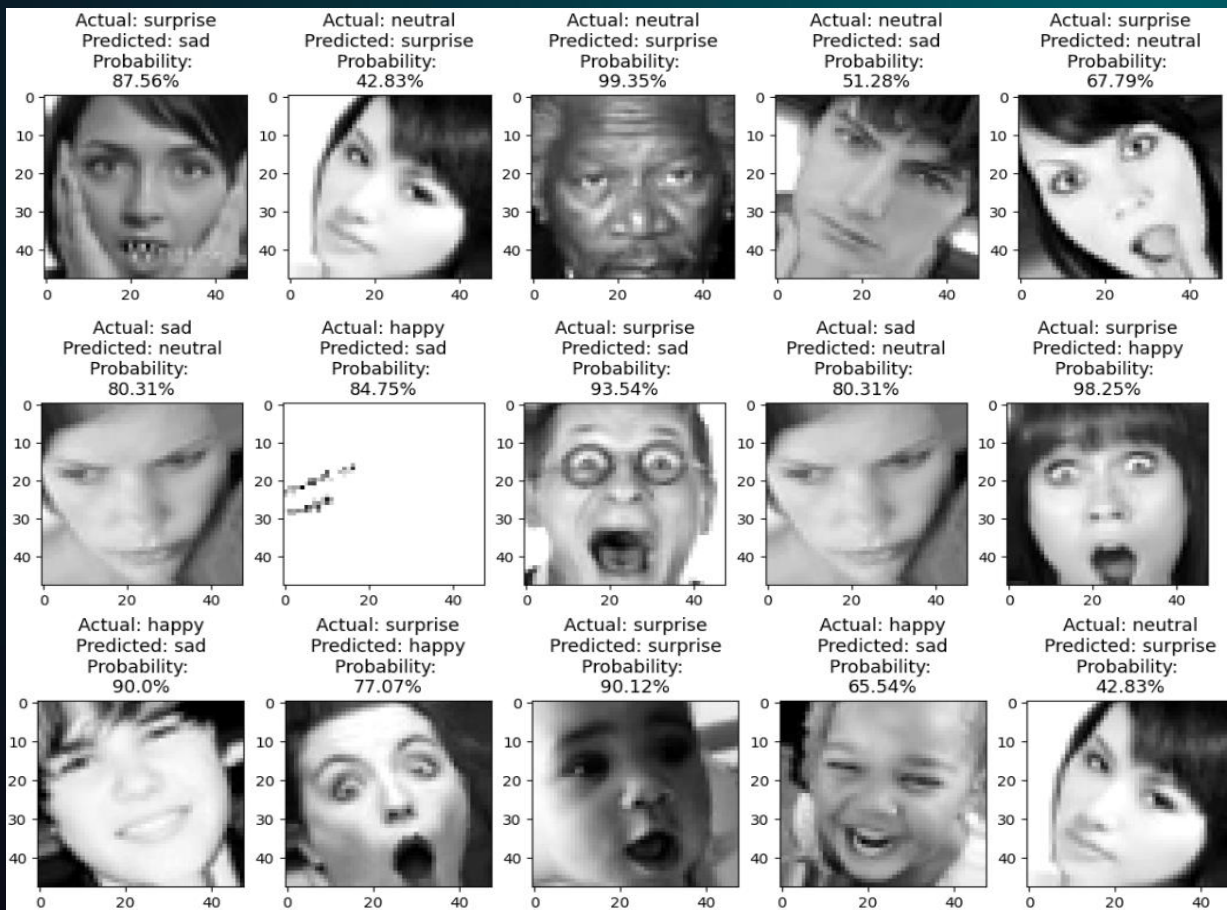


FINAL MODEL HEATMAP

PREDICTED RESULTS vs ACTUAL (Folder)



FINAL MODEL MISMATCH IMAGES



FINAL CNN MODEL ANALYSIS

	Parameters	Train Accuracy	Train Loss	Val Accuracy	Val Loss	Test Accuracy	Test Loss
Model 1: Grayscale 1	605,060	0.706	0.725	0.68	0.79	0.62	0.81
Model 2: Grayscale 2	389,604	0.796	0.496	0.715	0.74	0.68	0.68
Model 3: VGG16	14,714,688	0.72	0.687	0.675	0.815	0.72	0.71
Model 4: ResNet V2	42,658,176	0.27	1.38	0.36	1.37	0.25	1.40
Model 5: EfficientNet	8,769,374	0.27	1.38	0.37	1.35	0.25	1.41
Model 6: Grayscale 3	1,592,324	0.80	0.51	0.77	0.60	0.805	0.543
Model 7: Grayscale 4	2,973,700	0.83	0.44	0.77	0.65	0.77	0.60

CONCLUSIONS

- Less complex models work better with Facial_emotion_images
- Transfer learning models mediocre to poor on dataset



CONCLUSIONS

- How did the models come out Mr Bean?

